Wendong Li

| Toronto, Canada | (647) 233-7096 | levscaut@gmail.com | LinkedIn | GitHub

Experience

Microsoft **Toronto, Canada**
Software Engineer II Jul. 2025 – Present

Fabric IQ (AI agents over Power BI and Microsoft Fabric data)

- Built the evaluation platform for Fabric IQ agents across M365 Copilot, Copilot Cowork, and GitHub Copilot CLI: runners, versioned datasets, shared judging, and dashboards for engineers, scientists, and leadership.
- Drove cross-team adoption, onboarding partner teams' agent evaluations onto the platform with their own runners on shared scoring and dashboards; presented the tooling at an org-level execution review.
- Automated daily runs of the 352-query golden set against live MCP tools, with reliable parallel execution, judging, and lakehouse reporting; cut test-environment setup from days to minutes.
- Built a root-cause analysis agent with a custom harness and read-scoped tools over OpenTelemetry traces, recovering 97% of human-found root causes on a dogfooding benchmark; deployed in daily pipelines and the dashboard and packaged as agent skills, saving 2–3 hours of manual triage per experiment.
- Co-developed the golden set with an external vendor and drove its quality: benchmarked LLM judges, traced 84% of pass-to-fail flips to ambiguous rubrics, and fixed the rubrics with dev and PM, nearly halving failing queries.
- Implemented multi-turn and visual-answer evaluation, preserving per-turn tool calls and assertions and converting rendered charts and cards into image evidence for vision-capable judges.
- Shipped a regression alerting service for agent quality and latency with idempotent queue delivery; adopted by partner teams to monitor their daily evaluations.

Multimodal AI Functions (LLM transforms over DataFrames)

- Led multimodal AI Functions from idea to public preview in pandas and PySpark; showcased in the FabCon Atlanta 2026 opening keynote and a 350+ attendee session.
- Ran the private preview with 10+ clients through office hours and a cross-team bug bash; built the FabCon Vienna 2025 demo and helped an enterprise customer resolve context-window limits.
- Built OpenAI Responses API extraction for PDF, image, and lakehouse inputs with schema inference; replaced fuzzy matching with strict output schemas, cutting sentiment-analysis input tokens by ~35%.
- Red-teamed multimodal AI Functions for harmful content and prompt injection; hardened grammar-correction prompts, cutting attack success rate from 43% to 19% (most other functions: 0%).

Microsoft **Beijing, China**
Software Engineer Intern Feb. 2023 – Dec. 2023

FLAML and MLflow

- Benchmarked FLAML (AutoML) on Fabric against Azure ML and Databricks; integrated MLflow tracking, versioning, and serving; built a feature-engineering agent to automate data-science experiments.

Education

University of Toronto **Toronto, Canada**
M.Eng., Computer Engineering Jan. 2024 – Jun. 2025

Beijing University of Posts and Telecommunications **Beijing, China**
B.Eng., Software Engineering Sep. 2019 – Jun. 2023

Open Source

AutoGen **Open-source contributor**

- Contributed to the early design and implementation of AutoGen's RAG module, benchmarking vector databases and integrating ChromaDB for retrieval.
- Contributed Claude integration examples through three merged pull requests in April 2024, implementing a custom model client, system-prompt handling, and function calling. Model client | Tool calling

Technical Skills

Languages and Data: Python, SQL, JavaScript, Scala, pandas, PySpark, MLflow
Agents and Evaluation: MCP, tool calling, GitHub Copilot SDK, AutoGen, RAG, structured outputs, LLM-as-judge, multimodal evaluation, red-teaming
Platform: Microsoft Fabric, Azure OpenAI, Flask, SQLite, OpenTelemetry, Azure DevOps CI/CD